



OPERATING SYSTEMS

UE24CS242B

System calls for Process Management

Pavan A C

Department of Computer Science &
Engineering

Operating Systems

System calls for Process Management

Pavan A C

Department of Computer Science and Engineering

- Every process has a unique **process ID**, a non-negative integer
- The process ID is the only well-known identifier of a process that is always unique
- It is often used as a piece of other identifiers, to guarantee uniqueness.
- Although unique, process IDs are reused. As processes terminate, their IDs become candidates for reuse.
- Most UNIX systems implement algorithms to delay reuse, so

- An existing process can create a new one by calling the **fork** function

```
#include <unistd.h>

pid_t fork(void);
```
- The new process created by fork is called the child process
- Return value in the child is 0
- Return value in the parent is the process ID of the new child
- Return value is -1 on error
- The child is a copy of the parent. The child gets a copy of

- The `fork()` function creates the new process, just like `fork()`, without copying the address space of the parent into the child, as the child won't reference that address space; the child simply `#include <unistd.h>` `exit` (or `exit`) right after the `fork`

```
pid_t fork(void);
```

- The new process created by `fork` is called the child process
- Return value in the child is 0
- Return value in the parent is the process ID of the new child
- Return value is -1 on error
- `fork` guarantees that the child runs first, until the child

- A process can terminate normally in five ways
 1. Executing a return from the main function
 2. Calling the **exit** function.
 3. Calling the `_exit` or `_Exit` function.
 4. Executing a return from the start routine of the last thread in the process.
 5. Calling the `pthread_exit` function from the last thread in the process

A process can terminate abnormally in three ways

6. Calling `abort`
 7. When the process receives certain signals
 8. The last thread responds to a cancellation request
- Regardless of how a process terminates, the same code in the kernel is eventually executed.

- A process that calls wait or waitpid can
 - Block, if all of its children are still running
 - Return immediately with the termination status of a child, if a child has terminated and is waiting for its termination status to be fetched
 - Return immediately with an error, if it doesn't have any child processes
- If the process receives the SIGCHLD signal, wait will return immediately. If it is interrupted at any random point in time, it can block.

```
#include <sys/wait.h>
pid_t wait(int *statloc);
pid_t waitpid(pid_t pid, int *statloc, int options);
```

- waitid() allows a process to specify which children to wait for.
- Instead of encoding this information in a single argument combined with the process ID or process group ID, two separate arguments are used

```
#include <sys/wait.h>
```

```
int waitid(idtype_t idtype, id_t id, siginfo_t *infop, int options);
```

- Returns: 0 if OK, -1 on error

- When a process calls one of the exec functions, that process is completely replaced by the new program and the new program starts executing at its main function.
- The process ID does not change across an exec, because a new process is not created;

- exec() replaces the text, data, heap, and

stack and memory from disk.

```
#include <unistd.h>

int execl(const char *pathname, const char *arg0, ... /* (char *)0 */ );
int execlp(const char *pathname, char *const argv[]);
int execlx(const char *pathname, const char *arg0, ...
           /* (char *)0, char *const envp[] */ );
int execve(const char *pathname, char *const argv[], char *const envp[]);
int execlp(const char *filename, const char *arg0, ... /* (char *)0 */ );
int execlvp(const char *filename, char *const argv[]);
int fexecve(int fd, char *const argv[], char *const envp[]);
```

```
#include <unistd.h>
```

```
pid_t getpid(void);
```

```
pid_t getppid(void);
```

- **getpid()** returns the process ID (PID) of the calling process.
- **getppid()** returns the process ID of the parent of the calling process.
- This will be either the ID of the process that created this process using **fork()** or if that process has already terminated, the ID of the process to which this process has

- A race condition occurs when multiple processes are trying to do something with shared data and the final outcome depends on the order in which the processes run.
- The fork function is a lively breeding ground for race conditions, if the logic after the fork either explicitly or implicitly depends on whether the parent or child runs first after the fork.
- In general, which process runs first cannot be predicted
- ```
while (getppid() != 1)
 sleep(1);
```

its to wait for a child to terminate must call one of the wait functions.
- If a process wants to wait for its parent to terminate, a loop of the following form could be used:

- The problem with this type of loop called polling is that it wastes

- An I / O request is a request made by a running process to access an input or output device.
- Input devices include keyboard, mouse, and network interfaces, while output devices include display and printers.
- I / O operations are significantly slower than CPU operations, so they must be managed efficiently.
- The operating system acts as an intermediary between user processes and hardware devices.
- The I / O life cycle defines the sequence of steps involved from request initiation to completion.
- Proper I / O management improves system performance, responsiveness, and resource utilization.

## Life Cycle of an I/O Request – I/O Request Initiation and Kernel Processing

---

- A process initiates an I/O operation using system calls such as `read()`, `write()`, or `open()`.
- The system call causes a switch from user mode to kernel mode to ensure safe execution.
- The kernel validates the request parameters, file descriptors, and access permissions.
- The operating system identifies the target device and corresponding device driver.
- The I/O request is encapsulated into a kernel data structure.
- The request is placed into the I/O queue of the appropriate device.
- Depending on the type of I/O, the request may be handled as blocking or non-blocking.

## Life Cycle of an I/O Request – Device Driver and I/O Execution

---

- The device driver acts as a communication layer between the operating system and hardware.
- It converts generic kernel I/O requests into hardware-specific commands.
- The device controller is programmed to start the I/O operation.
- During I/O execution, the requesting process is typically blocked and moved to the waiting state.
- The operating system performs a context switch to allow other processes to execute.
- This overlap of CPU execution and I/O improves overall system throughput.

## Life Cycle of an I/O Request – I/O Completion and Process Resumption

---

- After completing the I/O operation, the device generates an interrupt signal.
- The interrupt notifies the CPU that the I/O request has finished.
- The operating system executes the Interrupt Service Routine (ISR).
- The ISR verifies the completion status and handles any errors.
- The blocked process is moved from the waiting state to the ready state.
- When the scheduler selects the process, it resumes execution.
- The system call completes and control is returned to user mode.



**THANK YOU**

---

**Pavan A C**

Department of Computer Science and  
Engineering  
**pavanac@pes.edu**