



OPERATING SYSTEMS

UE24CS242B

Mutex Locks and Semaphores

Pavan A C

Department of Computer Science &
Engineering

OPERATING SYSTEMS

- ~~Slides Credits for all the PPTs of this course~~
The slides/diagrams in this course are an adaptation, combination, and enhancement of material from the following resources and persons:

1. Slides of Operating System Concepts, Abraham Silberchatz, Peter Baer Galvin, Greg Gagne - 9th edition 2013 and some slides from 10th edition 2018
2. Some conceptual text and diagram from Operating Systems - Internals and Design Principles, William Stallings, 9th edition 2018
3. Some presentation transcripts from A. Frank – P. Weisberg
4. Some conceptual text from Operating System: Three Easy Pieces, Remzi Arpaci



OPERATING SYSTEMS

Mutex Locks and Semaphores

Pavan A C

Department of Computer Science & Engineering

OPERATING SYSTEMS

Mutex Locks



- ❑ Previous solutions i.e hardware and software solutions are complicated and generally inaccessible to application programmers
- ❑ OS designers build software tools to solve critical section problem
- ❑ Simplest is mutex lock
- ❑ Protect a critical section by first **acquire()** a lock then **release()** the lock
 - ❑ Boolean variable indicating if lock is available or not
- ❑ Calls to **acquire()** and **release()** must be atomic
 - ❑ Usually implemented via hardware atomic instructions
- ❑ But this solution requires **busy waiting**

Advantages of spinlocks

- ❖ no context switch is required when a process must wait on a lock.
- ❖ context switch may take considerable time.
- ❖ When locks are expected to be held for short times, spinlocks are useful
- ❖ These are employed on multiprocessor systems where one thread can “spin” on one processor while another thread performs its critical

Solution to Critical-section Problem Using Locks

```
while (TRUE) {  
    acquire lock  
    critical  
section  
    release lock  
remainder section  
}
```

- Implementation of acquire and release

1. While (LOCK==1);

2. LOCK = 1

3. Critical Section

4. LOCK = 0

Solution to Critical-section Problem Using Locks

```
while (TRUE) {  
    acquire lock  
    critical  
    section  
    release lock  
    remainder section  
}
```

- Implementation of acquire and release
- `acquire()` {
 while (!available)
 ; /* busy wait */
 available = false;
}
- `release()` {
 available = true;
}

acquire() and release()

- Solution to a critical section problem using mutex
- do {
 acquire lock
 critical section
 release lock
 remainder section
} while (true);

OPERATING SYSTEMS


Scenario 1

- Consider a library of an university with 10 rooms
- At a time one room can be used by only one student by informing the front desk for reading.
- Once he completes reading, he has to inform the front desk.
- Person at front desk knows how many rooms are available for use and how many are occupied, how many of them are waiting.
- Once the room is vacant, who will get the chance to occupy the room?



OPERATING SYSTEMS

Scenario 2



OPERATING SYSTEMS

Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- Semaphore **S** – integer variable
- Can only be accessed via two indivisible (atomic) operations
 - **wait()** and **signal()**
 - 4 Originally called **P()** and **V()**
- Definition of the **wait()** operation

```
wait(S) {
    while (S <= 0)
        ; // busy wait
    S- - ; }
```
- Definition of the **signal()** operation

```
signal(S) {
    S++ ; }
```

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
 - Same as a **mutex lock**
- Can solve various synchronization problems
- Consider P_1 and P_2 that require S_1 to happen before S_2

Create a semaphore “**synch**” initialized to 0

P1:

S_1 ;

signal (synch);

P2:

wait (synch);

- Must guarantee that no two processes can execute the **wait()** and **signal()** on the same semaphore at the same time
- Thus, the implementation becomes the critical section problem where the **wait** and **signal** code are placed in the critical section
 - Could now have **busy waiting** in critical section implementation
 - 4 But implementation code is short
 - 4 Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good

OPERATING SYSTEMS

Semaphore Implementation with no Busy waiting



- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- **typedef struct {**

OPERATING SYSTEMS

Semaphore Implementation with no Busy waiting (Cont.)

```
wai t ( semaphore * S) {  
    S->val ue- - ;  
    i f ( S->val ue < 0) {  
        add thi s process to S->l i st ;  
        bl ock() ;  
    }  
}  
  
si gna l ( semaphore * S) {  
    S->val ue++ ;  
    i f ( S->val ue >= 0) {  
        remove a process P from S->l i st ;  
        wake up( P) ;  
    } }  
}
```

- Incorrect use of semaphore operations:
 - signal (mutex) ... wait (mutex)
 - wait (mutex) ... wait (mutex)
 - Omitting of wait (mutex) or signal (mutex) (or both)
- Deadlock and starvation are possible.
- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let **S** and **Q** be two semaphores initialized to 1

P_0		P_1
wait(S);		wait(Q);
wait(Q);		wait(S);
...		...
signal(S);		signal(Q);

- **Starvation – indefinite blocking**
 - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - Solved via **priority-inheritance protocol**

□ When several tasks are waiting for the same critical resource, the task which is currently holding this critical resource is given the highest priority among all the tasks which are waiting for the same critical resource.

□ Now after the lower priority task having the critical resource is given the highest priority then the intermediate priority tasks can not preempt this task. This helps in avoiding priority inversion.

□ When the task which is given the highest priority

- Consider the scenario with three processes **P1**, **P2**, and **P3**.
 - **P1** has the highest priority, **P2** the next highest, and **P3** the lowest.
- Assume that **P3** is holding semaphore **S** and that **P1** is waiting for **S** to be released
- Assume that **P2** is assigned the CPU and preempts **P3**
 - **P3** is still holding semaphore **S**
 - **P1** is waiting for **S** to be released
- What has happened is that **P2** - a process with a lower priority than **P1** - has indirectly prevented **P3** from gaining access to the resource.



THANK YOU

Pavan A C

Department of Computer Science and
Engineering
pavanac@pes.edu